

# Introduction

# Background

- ▶ Purpose of the class: Concurrency!
- ▶ What is concurrency?

# Concurrency

- ▶ A *concurrent* program consists of several independent tasks
  - ▶ These can be unordered with respect to each other
  - ▶ Or they can be partially ordered with respect to each other
- ▶ That's not the same as parallelism
  - ▶ Parallelism = Performing several tasks simultaneously
- ▶ “Concurrency is about *dealing* with lots of things at once. Parallelism is about *doing* lots of things at once.”
  - ▶ <https://blog.golang.org/concurrency-is-not-parallelism>

# Concepts

- ▶ Three main concepts that give beginning programmers trouble...

# Iteration

- ▶ Life isn't one \_\_\_\_ thing after another. It's the same \_\_\_\_ thing over and over again.

(Edna St. Vincent Millay)

# Recursion

- ▶ To understand recursion, you must first understand recursion.  
(Unknown)

# Concurrency

- ▶ Major Premise: Sixty men can do a piece of work sixty times as quickly as one man.

Minor Premise: One man can dig a post-hole in sixty seconds

Conclusion: Sixty men can dig a post hole in one second.

(Ambrose Bierce)

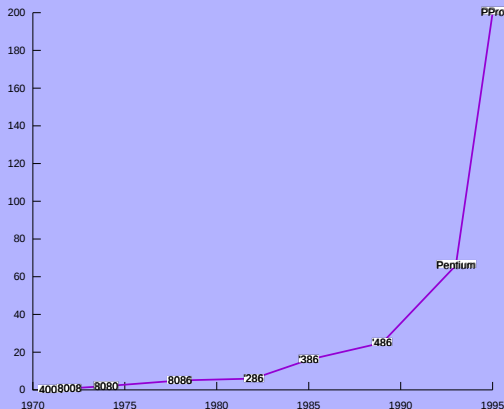
# Motivation

- ▶ Why concurrency?
- ▶ Two main reasons:
  - ▶ Speedup
  - ▶ Organization



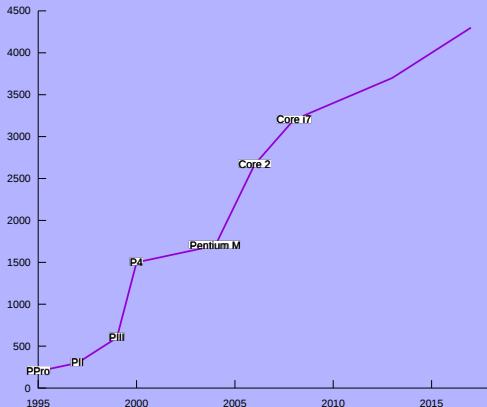
# Speed!

- ▶ Consider CPU speeds from early 70's through 1995: Exponential increase



# Speed!

- ▶ But then a funny thing happened...The rate of increase slowed down!



# Analysis

- ▶ Speeds are (slowly) hitting a brick wall
  - ▶ They creep upwards, but relative increases are getting smaller
- ▶ Limits of physics apply
  - ▶ Speed of light: 100 picoseconds (10GHz): 1.2 inches
  - ▶ Electrical charge: Move electrons on or off mosfet
    - ▶ There's a floor to how few electrons we can require (tunnelling)
  - ▶ SEU (cosmic rays) & transistor size

## The Past

- ▶ Is your program too slow?
  - ▶ Old joke: “You could spend a year rewriting it with a better algorithm or you could take a year off surfing.”
  - ▶ Either way, program will be fast enough next year!
- ▶ Not anymore...

# The Future

- ▶ Multiple cores are the trend
- ▶ As of late 2018:
  - ▶ AMD has announced a 48 core CPU (with 64 to follow)
  - ▶ Intel sells a 28 core unit (A bargain at \$9,990)
  - ▶ Intel has demonstrated a 72 *core* unit

## Make It Go Faster!

- ▶ Can't just wait a year!
  - ▶ Multiple cores require programmer to do the work
  - ▶ Break task into pieces; separate piece per core

# Complication

- ▶ The way the processor industry is going, is to add more and more cores, but nobody knows how to program those things.

(Steve Jobs, as reported by Peter Clarke [[https://www.eetimes.com/author.asp?section\\_id=36&doc\\_id=1266023](https://www.eetimes.com/author.asp?section_id=36&doc_id=1266023)])

## Challenge

- ▶ We need to use different programming techniques
- ▶ Concurrency demands a different tool set than nonconcurrent programming
- ▶ No industry consensus (yet) on the “best” way to approach the problem



# Measurements

- ▶ How do we know if we're making things better?
- ▶ Given two versions of same program:
  - ▶ Sequential one: Takes time

$$T_s$$

- ▶ Parallel one: Takes time

$$T_p$$

- ▶ Concept: Speedup:

$$S = \frac{T_s}{T_p}$$

- ▶ Ideally, if we throw  $n$  cores at problem, time taken is  $\frac{1}{n}$  original time
  - ▶ Thus, speedup =  $n$

## Limitation

- ▶ Any program can be divided into two parts:
  - ▶ Parallelizable part
  - ▶ *Inherently sequential* part
    - ▶ Must be done sequentially
    - ▶ Ex: Initialization, task decomposition, I/O, etc.
- ▶ Suppose program 10% inherently sequential and 90% parallelizable
- ▶ What can we expect?

## Theoretically

- ▶ Let  $T_s = 1$
- ▶ Suppose we have  $n$  CPU's
- ▶ Speedup:

$$S = \frac{T_s}{T_p} = \frac{1}{0.1 \cdot 1 + \frac{0.9 \cdot 1}{n}}$$

- ▶ Maximum possible speedup:

$$\lim_{n \rightarrow \infty} \frac{1}{0.1 \cdot 1 + \frac{0.9 \cdot 1}{n}} = \frac{1}{0.1} = 10$$

Hmmm...

- ▶ Let that sink in a moment
- ▶ We can *never* get more than 10x speedup!

## Moral

- ▶ We must work to reduce sequential part of code to absolute minimum if we want to get good performance
- ▶ Note that often parallel programming allows us to solve larger problems
  - ▶ So even if we don't see speedup, the fact that we can solve a larger problem → still a benefit.

# Distributed Systems

- ▶ What if we need more speedup than we can get on a desktop system?
- ▶ Ex: If we have a 16 core machine and every last bit of our code is parallelizable, we can get at most  $16\times$  speedup
- ▶ But what if we need to go faster?
- ▶ Answer: Distributed systems

# Distributed System

- ▶ What is a distributed system?
  - ▶ Several independent computers working together
  - ▶ Independent CPU's, separate RAM, usually independent storage (HD,SSD)
- ▶ Cloud computing
  - ▶ Many companies provide “rentable computing power”
    - ▶ Amazon AWS, Microsoft Azure, Digital Ocean, ...

# Concurrency

- ▶ Sometimes, we don't really care (so much) about speedup
  - ▶ We can use concurrent programming to *organize* our code
  - ▶ Here the goal is not so much speeding things up as helping us manage several activities simultaneously



# Concurrency

- ▶ Example: In a game program: What do we have going on at once?
  - ▶ Rendering
  - ▶ Audio
  - ▶ AI/Pathfinding
  - ▶ Network communication
  - ▶ I/O: Loading game world/textures/etc. in the background
- ▶ We need to manage these, even if we only have one core

# Topics

- ▶ What this class will cover:
  - ▶ How to create parallel/concurrent programs
  - ▶ Organization: How to organize these programs
  - ▶ Message passing: Communicating between tasks
  - ▶ Work distribution: How to divide up tasks

# Assignment

- ▶ Install either:
  - ▶ Visual Studio + C# toolchain
  - ▶ Visual Studio Code + C# module
- ▶ Run the obligatory Hello, World program:

```
1 using System;
2 class Hello{
3     public static void Main(string[] args){
4         Console.WriteLine("Hello, world");
5         Console.ReadLine();
6     }
7 }
```

- ▶ Turn in a screenshot of the program running on your computer

# Sources

- ▶ Intel Corporation:  
<http://www.intel.com/pressroom/kits/quickrefyr.htm> ,  
<http://www.intel.com/pressroom/kits/quickreffam.htm> ,  
<http://www.intel.com/pressroom/kits/quickrefyr.htm> ,  
<http://www.intel.com/content/www/us/en/history/history-intel-chips-timeline-poster.html?wapkw=processor+timeline>  
[http://ark.intel.com/products/93790/Intel-Xeon-Processor-E7-8890-v4-60M-Cache-2\\_20-GHz](http://ark.intel.com/products/93790/Intel-Xeon-Processor-E7-8890-v4-60M-Cache-2_20-GHz)
- ▶ Intel Corporation. The Evolution of a Revolution.  
<http://download.intel.com/pressroom/kits/IntelProcessorHistory.pdf>
- ▶ Intel Corporation. Intel Ark. <http://ark.intel.com/>
- ▶ AMD, inc.  
<http://products.amd.com/en-us/OpteronCPUDetail.aspx?id=814>
- ▶ Wikipedia. Transistor Count.
- ▶ <http://www.extremetech.com/extreme/171678-intel-unveils-72-core-x86-knights-landing-cpu-for-exascale-supercomputing>
- ▶ <http://www.amd.com/en-us/products/server/opteron/6000/6300#>
- ▶ [https://www.reddit.com/r/gamedev/comments/44fux4/multi\\_threading\\_in\\_game\\_development/](https://www.reddit.com/r/gamedev/comments/44fux4/multi_threading_in_game_development/)
- ▶ <https://www.extremetech.com/computing/245462-amds-new-32-core-naples-server-cpu-will-attack-intels-datacenter-monopoly>

Created using L<sup>A</sup>T<sub>E</sub>X.

Main font: Gentium Book Basic, by Victor Gaultney. See <http://software.sil.org/gentium/>

Monospace font: Source Code Pro, by Paul D. Hunt. See <https://fonts.google.com/specimen/Source+Code+Pro> and <http://sourceforge.net/adobe>

Icons by Ulisse Perusin, Steven Garrity, Lapo Calamandrei, Ryan Collier, Rodney Dawes, Andreas Nilsson, Tuomas Kuosmanen, Garrett LeSage, and Jakub Steiner. See <http://tango-project.org>