

Locks

Recall

- Recall our example from last time: The output was not 3000000

```
1 using System.Threading;
2 using System.IO;
3 using System;
4 class MainClass{
5     static int result = 0;
6     static void worker(){
7         for(int i = 0; i < 1000000; ++i)
8             result++;
9     }
10    public static void Main(string[] args){
11        Thread t1 = new Thread( () => { worker(); } );
12        Thread t2 = new Thread( () => { worker(); } );
13        Thread t3 = new Thread( () => { worker(); } );
14        t1.Start(); t2.Start(); t3.Start();
15        t1.Join(); t2.Join(); t3.Join();
16        Console.WriteLine("Foo: " + result);
17    }
18 }
```

Why?

- ▶ Recall: CPU doesn't necessarily do increment as atomic operation
 - ▶ Load
 - ▶ Modify
 - ▶ Store

Race Condition

- ▶ Race condition (or correctness hazard): Correctness of program depends on scheduling order of two or more threads
 - ▶ If thread A “wins the race”: One set of results
 - ▶ If thread B wins: Different results
- ▶ All programs with race conditions are erroneous

Hazards

- ▶ When do they occur?
 - ▶ When contention for shared resource
 - ▶ Locals are never shared
 - ▶ So: If program only uses locals: No contention
- ▶ Note: Static instance variables are inherently shared
- ▶ Note: If you pass same reference to two threads: Shared resource

Globals

- ▶ Most useful programs must use global resources
 - ▶ Shared data (ex: Large array – too costly to duplicate)
 - ▶ Common I/O device (terminal)
 - ▶ Common disk file
 - ▶ We often use globals to communicate between threads
- ▶ So it's hard to avoid potential for race conditions

Visibility

- ▶ Another concern: Visibility
- ▶ Suppose we have code like so:

```
1 static bool flag1=false;
2 static bool flag2=false;
```

```
1 void func1(){
2     doSomething();
3     flag1=true;
4     while(flag2 == false )
5         ;
6     doSomethingElse()
7 }
```

```
1 void func2(){
2     while(flag1 == false )
3         ;
4     doStuff();
5     flag2=true;
6     finishUp();
7 }
```

Problems

- ▶ First problem: Busy waiting
 - ▶ Wastes CPU
 - ▶ Creates heat
 - ▶ Burns battery life
- ▶ Second problem: No guarantee that this code does what we want
 - ▶ As long as code is correct from single thread's viewpoint, instructions can be reordered
 - ▶ So: Legal for compiler (or hardware) to reorder the write of flag1 until after the while loop
- ▶ If we replaced assignments with:
 Interlocked.Increment(flag1)
or
 Interlocked.Increment(flag2)
- ▶ Then: Code works. Why?

Reasoning

- ▶ Concept: Interlocked operations impose *sequential consistency*
- ▶ What's that?

Definition

- ▶ Define a relation: *happens-before*
- ▶ If statement X happens-before statement Y, it means... What you'd think!
 - ▶ Results of X are visible when Y begins
 - ▶ Written $X \rightarrow Y$
- ▶ If we don't know that $X \rightarrow Y$ and we don't know that $Y \rightarrow X$ then we say that X and Y are *concurrent*
 - ▶ Written $X \parallel Y$

Consider

- ▶ Consider previous code
- ▶ Which statement(s) happen before which other statements?

Result

- ▶ Here's all we can say:
 - ▶ $a \rightarrow b \rightarrow c \rightarrow d \rightarrow e$
 - ▶ $f \rightarrow g \rightarrow h \rightarrow i \rightarrow j$
- ▶ Hmm. Notice we can say nothing with regard to inter-thread operations
- ▶ The two threads are entirely *concurrent*
 - ▶ $a \parallel f, a \parallel g, a \parallel h, b \parallel f, b \parallel g$, etc.
- ▶ This is different when interlocked operations are used...

Interlocked

- ▶ If we use interlocked operations, this creates a synchronization point
- ▶ Suppose threads A and B access *the same atomic variable x*
 - ▶ Thread A reads x at statement α
 - ▶ Thread B writes x at statement β
 - ▶ Then exactly one of these will be true:
 $\alpha \rightarrow \beta$
 $\beta \rightarrow \alpha$
- ▶ This is the crucial inter-thread tie that we need!

New Code

```
1 static int flag1=0;
2 static int flag2=0;
```

```
1 void func1(){
2     doSomething();
3     //(a)
4     Interlocked.Increment(flag1);
5     //(b)
6     while( Interlocked.Add(flag2
7         ,0) == 0 )
8         ;
9     //(c)
10    doSomethingElse()
11 }
```

```
1 void func2(){
2     while(Interlocked.Add(flag1
3         ,0) == 0 ) //(e)
4         ;
5     doStuff();
6     //(f)
7     Interlocked.Increment(flag2);
8     //(g)
9     finishUp();
10    //(h)
11 }
```

Relationships

- ▶ $a \rightarrow b \rightarrow c \rightarrow d$
- ▶ $e \rightarrow f \rightarrow g \rightarrow h$
- ▶ We know that $b \rightarrow f$
 - ▶ Why? Because:
 - ▶ $e \rightarrow f$
 - ▶ We get past e iff flag1 is nonzero
 - ▶ b is the only place we set flag1 to nonzero
 - ▶ Thus, $(b||e) \rightarrow f$
- ▶ Likewise, we know $f \rightarrow d$
 - ▶ Why? Because $f \rightarrow g \rightarrow d$
 - ▶ Do you see why?

Interlocked

- ▶ Interlocked solves some visibility and ordering problems
- ▶ But: Important drawbacks:
 - ▶ Don't help with busy waiting
 - ▶ Difficult to reason about program correctness
 - ▶ Don't help with ensuring several interdependent variables are kept consistent

Mutex

- ▶ Mutex = MUTual EXclusion
 - ▶ Also called a lock
- ▶ A mutex is essentially a boolean with two operations: lock and unlock

Mutex

- ▶ Lock: Pseudocode:

```
1 while( locked == true )  
2     releaseCpu();  
3 locked=true;
```

- ▶ What's not shown here is that all operations are atomic
 - ▶ Locked can't be changed after leaving the loop but before setting locked=true
 - ▶ We normally can't write code like that ourselves; needs OS/runtime support

Mutex

► Unlock: Pseudocode:

```
1 locked=false;
```

Vocabulary

- ▶ If thread A successfully calls mutex lock (i.e., when lock() returns), we say A has *acquired* the mutex (or "it has locked the lock")
- ▶ If A calls unlock(), we say it has *released* the mutex (or "unlocked the lock")

Vocabulary

- ▶ What happens if A acquires the lock and then B calls lock()?
 - ▶ OS takes B off the CPU until A releases the mutex
 - ▶ We say B is *blocked* on the mutex (or B is "waiting for the lock")
 - ▶ The CPU is free for other threads while B is blocked
 - ▶ So we *are not* busy waiting

Syntax

- ▶ In C#, to perform mutex operation: We first create an object and ensure all threads can see it:

`object myLock = new object()`

- ▶ Then, we write:

```
1 lock( myLock ){  
2     ...  
3 }
```

- ▶ Lock will be held for duration of brace-block

Sequential Consistency

- ▶ Mutexes (mutices?) provide some sequential consistency guarantees as well

```
1 static object myLock = new object();
```

```
1 void func1(){  
2     //(a)  
3     lock(myLock){  
4         //(b)  
5         doSomething();  
6         //(c)  
7     }  
8     //(d)  
9 }
```

```
1 void func2(){  
2     //(e)  
3     lock(myLock){  
4         //(f)  
5         doSomething();  
6         //(g)  
7     }  
8     //(h)  
9 }
```

Ordering

- ▶ Trivially, we know:

$a \rightarrow b \rightarrow c \rightarrow d$

$e \rightarrow f \rightarrow g \rightarrow h$

- ▶ Since we are locking on same lock, we know either:

$d \rightarrow f$ or

$h \rightarrow b$

- ▶ So we can reason about inter-thread dependencies this way

Rule of Thumb

- ▶ To make life easier, we have some common patterns we use with mutexes
- ▶ Pattern: If you have a shared variable: you guard it with a mutex
 - ▶ Note: Every thread must use the same mutex to get any useful synchronization
 - ▶ In practice, that means your mutex object will almost always be static

Incorrect

```
1 void func1(){  
2     object M = new object();  
3     lock(M){  
4         ...  
5     }  
6 }
```

- ▶ M is totally useless!

Example

- ▶ Suppose we have a network server
- ▶ Remote machines contact it, request resources, get data back
- ▶ Suppose it takes time to compute response, so we cache last request to save time

Example

- ▶ Example non-threaded server implementation:

```
1 class Server{
2     static DataItem lastItem;
3     static string lastIdentifier;
4     DataItem handleRequest(string identifier){
5         if( identifier == lastIdentifier )
6             return lastItem;
7         else{
8             lastItem = loadDataFromDisk(identifier);
9             lastIdentifier = identifier;
10            return lastItem;
11        }
12    }
13 }
```

- ▶ Clearly this would not be multi-thread-safe. Why not?

Solution

```
1 class Server{
2     static DataItem lastItem;
3     static string lastIdentifier;
4     static object M = new object();
5     DataItem handleRequest(string identifier){
6         lock(M){
7             if( identifier == lastIdentifier )
8                 return lastItem;
9             else{
10                 lastItem = loadDataFromDisk(identifier);
11                 lastIdentifier = identifier;
12                 return lastItem;
13             }
14         }
15     }
16 }
```

Extension

- ▶ Now, we decide to extend it: Save the last five requests in the cache. The non-threaded code:

```
1 class Server{
2     static DataItem[] lastItem = new DataItem[5];
3     static string[] lastIdentifier = new string[5];
4     DataItem handleRequest(string identifier){
5         for(int i=0;i<5;++i){
6             if( identifier[i]== lastIdentifier[i] )
7                 return lastItem[i];
8         }
9         //FIFO replacement strategy
10        for(int i=0;i<4;++i){
11            lastItem[i] = lastItem[i+1];
12            lastIdentifier[i] = lastIdentifier[i+1];
13        }
14        lastItem[4] = loadDataFromDisk(identifier);
15        lastIdentifier[4] = identifier;
16        return lastItem[4];
17    }
18 }
```

Threaded

- Sprinkle some magic locking pixie dust around...Broken!

```
1 class Server{
2     static DataItem[] lastItem = new DataItem[5];
3     static string[] lastIdentifier = new string[5];
4     object M = new object();
5     DataItem handleRequest(string identifier){
6         lock(M){
7             for(int i=0;i<5;++i){
8                 if( identifier[i]== lastIdentifier[i] ) return lastItem[i];
9             }
10            //FIFO replacement strategy
11            for(int i=0;i<4;++i){
12                lastItem[i] = lastItem[i+1];
13                lastIdentifier[i] = lastIdentifier[i+1];
14            }
15            lastItem[4] = loadDataFromDisk(identifier);
16            lastIdentifier[4] = identifier;
17            return lastItem[4];
18        }
19    }
20 }
```

Problem

- ▶ If one thread is computing, it blocks all other threads from doing any work
 - ▶ Even if they could go immediately!
- ▶ How to fix?

Solution?

► How about this?

```
1 class Server{
2     static DataItem[] lastItem = new DataItem[5];
3     static string[] lastIdentifier = new string[5];
4     object M = new object();
5     DataItem handleRequest(string identifier){
6         for(int i=0;i<5;++i){
7             if( identifier[i]== lastIdentifier[i] ) return lastItem[i];
8         }
9         lock(M){
10             //FIFO replacement strategy
11             for(int i=0;i<4;++i){
12                 lastItem[i] = lastItem[i+1];
13                 lastIdentifier[i] = lastIdentifier[i+1];
14             }
15             lastItem[4] = loadDataFromDisk(identifier);
16             lastIdentifier[4] = identifier;
17             return lastItem[4];
18         }
19     }
20 }
```

Nope

- ▶ We're accessing shared data without holding the lock!

Fixed? (Nope!)

► Is this good? Why not?

```
1 class Server{
2     static DataItem[] lastItem = new DataItem[5];
3     static string[] lastIdentifier = new string[5];
4     object M = new object();
5     DataItem handleRequest(string identifier){
6         lock(M){
7             for(int i=0;i<5;++i)
8                 if( identifier[i] == lastIdentifier[i] ) return lastItem[i];
9             for(int i=0;i<4;++i){
10                 lastItem[i] = lastItem[i+1];
11                 lastIdentifier[i] = lastIdentifier[i+1];
12             }
13         }
14         var tmp = loadDataFromDisk(identifier);
15         lock(M){
16             lastItem[4] = tmp;
17             lastIdentifier[4] = identifier;
18         }
19         return lastItem[4];
```

Problem

- ▶ Race condition after releasing lock but before returning `lastIdentifier[4]`

Solution (Finally!)

► Finally!

```
1 class Server{
2     static DataItem[] lastItem = new DataItem[5];
3     static string[] lastIdentifier = new string[5];
4     object M = new object();
5     DataItem handleRequest(string identifier){
6         lock(M){
7             for(int i=0;i<5;++i){
8                 if( identifier[i]== lastIdentifier[i] )
9                     return lastItem[i];
10            }
11        }
12        var tmp = loadDataFromDisk(identifier);
13        lock(M){
14            //FIFO replacement strategy
15            for(int i=0;i<4;++i){
16                lastItem[i] = lastItem[i+1];
17                lastIdentifier[i] = lastIdentifier[i+1];
18            }
19            lastItem[4] = tmp;
20            lastIdentifier[4] = identifier;
21            return lastItem[4];
22        }
23    }
24 }
```

Deadlocks

- ▶ Deadlocks are one of the main hazards when working with mutexes
- ▶ Suppose we have two global variables, each protected by its own lock

```
1 static object xLock = new object();
2 static object yLock = new object();
3 static int x,y;
```

```
1 void func1(){
2     lock(xLock){
3         if( x > 5 ){
4             lock(yLock){
5                 y++;
6             }
7         }
8     }
9 }
```

```
1 void func2(){
2     lock(yLock){
3         if( y > 100 ){
4             lock(xLock){
5                 x++;
6             }
7         }
8     }
9 }
```

Problem

- ▶ classic x-y / y-x locking pattern
- ▶ This puts you on the express train to Deadlockville

Solutions

- ▶ One solution: Number all locks
- ▶ Then: Ensure locks are acquired in increasing numerical order
 - ▶ Ex: Let xLock be 0 and yLock be 1
 - ▶ If you want both locks, you must grab xLock first
- ▶ Problem: Second thread doesn't know right away if it needs xLock
 - ▶ Maybe y is 42
 - ▶ Lower resource utilization
 - ▶ Need to grab lots of locks "just in case" we need them later
 - ▶ Hard to remember the number of each lock

Assignment

- ▶ Extend the previous lab:
 - ▶ Every half second, print the status of each downloaded item: Either “Complete” or “In progress” or “Error”
- ▶ When all items are downloaded (or have errored out), exit.
- ▶ Turn in your CS files only
- ▶ More details follow...

Assignment

- ▶ Your code must have no race conditions, no possibility of incorrect output, and no visibility hazards
- ▶ Your program must not crash; make sure to handle all exceptions
- ▶ Ex: Your display might look like this:

```
1 http://www.example.com: In progress
2 http://www.example.org/foo/bar: Complete
3 http://www.example.net/abc.txt: In progress
4 http://www.example.org/def.gif: Error
5 -----
6 http://www.example.com: In progress
7 http://www.example.org/foo/bar: Complete
8 http://www.example.net/abc.txt: Complete
9 http://www.example.org/def.gif: Error
10 -----
11 http://www.example.com: Complete
12 http://www.example.org/foo/bar: Complete
13 http://www.example.net/abc.txt: Complete
14 http://www.example.org/def.gif: Error
```

Sources

- ▶ B. Goetz et al. Java Concurrency in Practice. Addison-Wesley.
- ▶ M. Herlihy & N. Shavit. The Art of Multiprocessor Programming. Morgan Kaufmann Publishers.
- ▶ David Beazley & Brian K. Jones. Python Cookbook. O'Reilly Media.
- ▶ Varun. C++ 11 Multithreading Part 5: Using Mutex to Fix Race Conditions. <http://thispointer.com/c11-multithreading-part-5-using-mutex-to-fix-race-conditions/>
- ▶ <http://www.ecma-international.org/publications/files/ECMA-ST/ECMA-335.pdf>

Created using L^AT_EX.

Main font: Gentium Book Basic, by Victor Gaultney. See <http://software.sil.org/gentium/>

Monospace font: Source Code Pro, by Paul D. Hunt. See <https://fonts.google.com/specimen/Source+Code+Pro> and <http://sourceforge.net/adobe>

Icons by Ulisse Perusin, Steven Garrity, Lapo Calamandrei, Ryan Collier, Rodney Dawes, Andreas Nilsson, Tuomas Kuosmanen, Garrett LeSage, and Jakub Steiner. See <http://tango-project.org>