

Semaphores

Motivation

- ▶ Mutexes: Simple, but some problems:
- ▶ Only two states: Locked and Unlocked
- ▶ Some problems are hard to solve with mutexes

Latch

- ▶ Scenario: We have several threads that are created at different times
- ▶ Want all to start processing (nearly) simultaneously
 - ▶ Ex: Multiplayer game
 - ▶ Start one thread as each player connects
 - ▶ Don't begin processing until all players ready
- ▶ How to structure?

Latch

► Attempt 1:

```
1 //global
2 static int numLeft = 5;
3 void threadFunc(){
4     while(numLeft > 0 ){
5     }
6     ...
7 }
8 void main(){
9     Thread t1 = new Thread( () => { threadFunc(); } );
10    numleft--;
11    Thread t2 = new Thread( () => { threadFunc(); } );
12    numleft--;
13    Thread t3 = new Thread( () => { threadFunc(); } );
14    numleft--;
15    ...etc...
16 }
```

► Why is this very wrong?

Analysis

- ▶ This is incorrect
- ▶ Visibility: Happens-before
- ▶ We need to use interlocked

Latch

► Attempt 2:

```
1 //global
2 static int numLeft = 5;
3 void threadFunc(){
4     while(Interlocked.Add(numLeft,0) ){
5     }
6     ...
7 }
8 void main(){
9     Thread t1 = new Thread( () => { threadFunc(); } );
10    Interlocked.Add(numLeft,-1);
11    Thread t2 = new Thread( () => { threadFunc(); } );
12    Interlocked.Add(numLeft,-1);
13    Thread t3 = new Thread( () => { threadFunc(); } );
14    Interlocked.Add(numLeft,-1);
15    ...etc...
16 }
```

► Still bad. Why?

Analysis

- ▶ Wastes lots of CPU time
 - ▶ CPU constantly checking variable
 - ▶ Battery life diminished
 - ▶ Other tasks become sluggish
 - ▶ Heat

Attempt 3

► Third attempt:

```
1 //global
2 static int numLeft = 5;
3 void threadFunc(){
4     while(Interlocked.Add(numLeft,0) ){
5         Thread.Sleep(1);    //sleep 1 msec
6     }
7     ...
8 }
9 void main(){
10     Thread t1 = new Thread( () => { threadFunc(); } );
11     Interlocked.Add(numLeft,-1);
12     Thread t2 = new Thread( () => { threadFunc(); } );
13     Interlocked.Add(numLeft,-1);
14     Thread t3 = new Thread( () => { threadFunc(); } );
15     Interlocked.Add(numLeft,-1);
16     ...etc...
17 }
```


Analysis

- ▶ A little better
- ▶ But some delay once all threads ready
- ▶ CPU must still “wake up” every so often
 - ▶ Still costs battery life
- ▶ We need some additional OS level facilities

Semaphore

- ▶ An integer with two operations
 - ▶ Acquire
 - ▶ Release

Naming

- ▶ Semaphores were invented by Edsger Dijkstra
- ▶ Several different names for operations
 - ▶ Acquire = down = wait = P
 - ▶ *Proberen* or *passeren*: To try or to pass
 - ▶ Release = up = signal = post = V
 - ▶ *Verhogen* or *verlaten* or *vrijgeven*: To increase or to leave or to release

Operations

- ▶ Can think of semaphore as collection of “permits”
- ▶ Upon acquire
 - ▶ If a permit is available, take it; else, wait (block).
- ▶ Upon release:
 - ▶ Give back a permit
 - ▶ Might immediately be taken up by waiting process, if any

Problem

- ▶ Create:
static Semaphore S = new Semaphore(initialValue, maxValue)
- ▶ Acquire:
S.WaitOne()
- ▶ Release:
S.Release()

Note

- ▶ No guarantee about wakeup order
- ▶ Just because task A wait's before task B doesn't mean that A gets awakened before B

Latch

- Now we can solve our latch problem:

```
1 //global
2 static Semaphore latch = (0,5);
3 void threadFunc(){
4     latch.WaitOne();
5     ...
6 }
7 void main(){
8     Thread t1 = new Thread( () => { threadFunc(); } );
9     Thread t2 = new Thread( () => { threadFunc(); } );
10    Thread t3 = new Thread( () => { threadFunc(); } );
11    ...etc...
12    for(int i=0;i<5;++i)
13        latch.Release();
14    ...
15 }
```

Rendezvous

- ▶ Want to ensure A has finished foo and bar before B executes boom and bash AND B finishes bam before baz can go
- ▶ Notation: $((\text{foo} \rightarrow \text{bar}) \parallel \text{bam}) \rightarrow (\text{baz} \parallel (\text{boom} \rightarrow \text{bash}))$

```
1 void A(){  
2     foo();  
3     bar();  
4     baz();  
5 }
```

```
1 void func2(){  
2     bam();  
3     boom();  
4     bash();  
5 }
```


Notation

- ▶ Recall: Our notation...
- ▶ If we write $x \parallel y$: Means x and y are *concurrent*
 - ▶ No assertion of relative order
 - ▶ Might happen truly in parallel on multicore machine
 - ▶ Or sequentially on single core machine
- ▶ If we write $x \rightarrow y$: Denotes *happens before*
 - ▶ x must complete before y starts

Example

- ▶ Suppose we have one thread:

foo();

bar();

baz()

- ▶ We must have $\text{foo} \rightarrow \text{bar} \rightarrow \text{baz}$

- ▶ In another thread:

bam()

boom()

bash()

- ▶ We must have $\text{bam} \rightarrow \text{boom} \rightarrow \text{bash}$

- ▶ Put it together: $(\text{foo} \rightarrow \text{bar} \rightarrow \text{baz}) \parallel (\text{bam} \rightarrow \text{boom} \rightarrow \text{bash})$

First Attempt

- ▶ Not quite right. Why not?

```
1 static Semaphore S = new Semaphore(0,Int32.MaxValue);
```

```
1 foo();  
2 bar();  
3 S.Release();  
4 baz();
```

```
1 bam();  
2 S.WaitOne();  
3 boom();  
4 bash();
```

Problem

- ▶ baz can complete before bam
 - ▶ Can you see how?

Notation

- ▶ We know:
 - ▶ $\text{foo} \rightarrow \text{bar} \rightarrow \text{baz}$
 - ▶ $\text{bar} \rightarrow \text{boom}$
 - ▶ Transitive: So $\text{foo} \rightarrow \text{boom}$
 - ▶ And $\text{boom} \rightarrow \text{bash}$
- ▶ But observe: $\text{baz} \parallel \text{bam}$
 - ▶ Diagram out: Easier to see...

Rendezvous

► Need two semaphores

```
1 static Semaphore S1 = new Semaphore(0,Int32.MaxValue);  
2 static Semaphore S2 = new Semaphore(0,Int32.MaxValue);
```

```
1 foo();  
2 bar();  
3 S1.Release();  
4 S2.WaitOne();  
5 baz();
```

```
1 bam();  
2 S2.Release();  
3 S1.WaitOne();  
4 boom();  
5 bash();
```

Deadlock

- ▶ Note that order is important
- ▶ If we reverse the Wait/Release: We get deadlock!

Multiplex

- ▶ Recall how we had mutexes previously
 - ▶ Kind of like binary semaphore: Semaphore with value 0 or 1
- ▶ Can we extend our mutex to a *multiplex*: Have a function `foo()` which allows no more than 10 threads to execute it at once?
 - ▶ Kind of a mirror image of latch problem from earlier

Multiplex

- ▶ Define shared variable:
`static Semaphore S = new Semaphore(10,10);`
- ▶ Function foo:
`void foo(){
 S.WaitOne();

 ...
 S.Release();
}`
- ▶ There's a potential problem here...

Problem

- ▶ What if something in foo (or one of its called functions) throws an exception?
 - ▶ Caller of foo might catch it
 - ▶ But foo will never increment the semaphore!
- ▶ Now our program is permanently broken

Solution

- ▶ In concurrency, this sort of problem appears frequently
- ▶ To solve:

```
1 void foo(){  
2     S.WaitOne();  
3     try{  
4         ...code...  
5     }  
6     finally{  
7         S.Release();  
8     }  
9 }
```

Assignment

- ▶ As explained in class.
- ▶ You are only allowed to change the Smoker.cs file, the Agent.cs file, and the Globals.cs file.

Sources

- ▶ B. Goetz et al. *Java Concurrency in Practice*. Addison Wesley.
- ▶ A. Downey. *The Little Book of Semaphores*.
<http://greenteapress.com/>
- ▶ M. Herlihy and N. Shvavit. *The Art of Multiprocessor Programming*
- ▶ Semaphore - Everything2.com
<http://everything2.com/title/Semaphore>
- ▶ Microsoft Corp. .NET framework documentation

Created using L^AT_EX.

Main font: Gentium Book Basic, by Victor Gaultney. See <http://software.sil.org/gentium/>

Monospace font: Source Code Pro, by Paul D. Hunt. See <https://fonts.google.com/specimen/Source+Code+Pro> and <http://sourceforge.net/adobe>

Icons by Ulisse Perusin, Steven Garrity, Lapo Calamandrei, Ryan Collier, Rodney Dawes, Andreas Nilsson, Tuomas Kuosmanen, Garrett LeSage, and Jakub Steiner. See <http://tango-project.org>