

Fair Semaphores

Motivation

- ▶ Recall how we discussed locks (mutexes)
- ▶ And we've seen semaphores
 - ▶ We can use semaphores like locks, too: Just restrict to values 0,1
- ▶ And we've seen the latch pattern

Latch

- ▶ Basic idea (pseudocode):
- ▶ Create global/static variable:
`S=new Semaphore(0,numThreads)`
- ▶ Main does this:
Start threads
`for(i=0;i<numThreads;++i)`
 `S.Release();`
- ▶ Threads do this:
 `S.acquire()`
 ...do stuff...

Question

- ▶ What if we don't want to have the main thread involved in the release?
- ▶ Just want threads to manage it themselves

Barrier

- ▶ Concept: The *barrier*
- ▶ All threads must get to some point in code, then all may continue
- ▶ Ex: Suppose we have a global object:
Barrier b
- ▶ And all threads are executing these statements:
 1. foo()
 2. b.wait()
 3. bar()
- ▶ We want to guarantee that $(A1 \parallel B1 \parallel C1) \rightarrow (A3 \parallel B3 \parallel C3)$
- ▶ How to code Barrier?

Barrier

► Pattern: Count + lock + semaphore

```
1 class Barrier{
2     private int max, num=0;
3     private object L = new object();
4     private Semaphore S;
5     public Barrier(int max){
6         this.max=max;
7         this.S = new Semaphore(0,max);
8     }
9     public void wait(){
10         lock(L){
11             num++;
12             if( num == max ){
13                 for(int i=0;i<max-1;++i)
14                     S.Release();
15             } else
16                 S.WaitOne();
17         }
18     }
19 }
```

Problem

- ▶ This is almost correct... Why does it fail?

Problem

- ▶ S.WaitOne while holding lock
- ▶ Deadlock!

Solution

- We need to use explicit locking pattern

```
1 class Barrier{
2     private int max, num=0;
3     private Semaphore S, L = new Semaphore(1,1);
4     public Barrier(int max){
5         this.max=max;
6         this.S = new Semaphore(0,max);
7     }
8     public void wait(){
9         L.WaitOne();
10        num++;
11        if( num == max ){
12            for(int i=0;i<max-1;++i)
13                S.Release();
14            L.Release();
15        } else {
16            L.Release();
17            S.WaitOne();
18        }
19    }
20 }
```

Reusable Barrier

- ▶ What if we can hit barrier several times?
- ▶ Ex: We have threads working like so:

```
while(true){  
    foo();  
    b.wait();  
    bar();  
}
```

- ▶ Desired semantics: For all threads t : When t calls `bar()` for the n th time, every other thread has called `bar` at least $n-1$ times
- ▶ Why doesn't our barrier from the previous slide work?

Problem

- ▶ It doesn't reset num

Solution?

► What if we re-code it like so:

```
1 class Barrier{
2     private int max, num=0;
3     private Semaphore S, L = new Semaphore(1,1);
4     public Barrier(int max){
5         this.max=max;
6         this.S = new Semaphore(0,max);
7     }
8     public void wait(){
9         L.WaitOne();
10        num++;
11        if( num == max ){
12            for(int i=0;i<max-1;++i)
13                S.Release();
14            num=0;
15            L.Release();
16        } else {
17            L.Release();
18            S.WaitOne();
19        }
20    }
```

Nope

- ▶ No dice. Why does this fail?

Problem

- ▶ Don't peek! See if you can figure it out on your own...

Problem

- ▶ Suppose $\text{max}=4$ and we have threads A,B,C call wait
- ▶ Suppose A and B get to S.WaitOne, but C is preempted between L.Release and S.WaitOne
- ▶ Thread D appears. It acquires L, sets num to 4, performs three Release's, sets num to 0, and releases L
- ▶ Thread A is awakened, leaves wait(), quickly does bar() and foo(), calls wait() again, gets L, sets num to 1, releases L, and then does S.WaitOne()
- ▶ What happens?

Problem

- ▶ A gets the Release that was meant for C!
- ▶ Unlikely? Maybe...
- ▶ But under pathological load, it can happen
- ▶ These are the worst bugs to find and fix!

Turnstile

- ▶ We'll go back and recode our initial solution using a *turnstile*
- ▶ This doesn't fix the cyclic barrier problem (yet), but it's a step in the right direction
- ▶ Turnstile pattern = Acquire semaphore immediately followed by release of that semaphore

Turnstile

► Code:

```
1 class Barrier{
2     private int max, num=0;
3     private object L = new object();
4     private Semaphore S = new Semaphore(0,1);
5     public Barrier(int max){
6         this.max=max;
7     }
8     public void wait(){
9         lock(L){
10             num++;
11             if( num == max )
12                 S.Release();
13         }
14         S.WaitOne();
15         S.Release();
16     }
17 }
```

Cyclic

- ▶ Now, we add a new function: `reset()`
- ▶ Idea: Each thread will have this format:

```
while(true){  
    foo();  
    b.wait();  
    bar();  
    b.reset()  
}
```
- ▶ After all threads have called `reset`: Barrier is “ready” again

Barrier

► We add the function:

```
1 class Barrier{
2     private int max, num=0;
3     private object L = new object();
4     private Semaphore S = new Semaphore(0,1);
5     public Barrier(int max){
6         this.max=max;
7     }
8     public void wait(){
9         lock(L){
10             num++;
11             if( num == max )
12                 S.Release();
13         }
14         S.WaitOne();
15         S.Release();
16     }
17     public void reset(){
18         lock(L){
19             //if I'm the first thread to call reset(),
20             //close the turnstile again
21             if( num == max )
22                 S.WaitOne();
23             num--;
24         }
25     }
26 }
```

Barrier


- ▶ It's broken. Why?

Problem

- ▶ Suppose four threads
- ▶ A,B,C blocked at S.WaitOne inside wait()
- ▶ D arrives. Does S.Release
- ▶ Exactly one of A,B,C will awaken from WaitOne
- ▶ Meanwhile, D is making its way toward the WaitOne
- ▶ Suppose A calls S.Release
- ▶ This will awaken exactly one of B, C, D (if D has gotten to WaitOne by now)
- ▶ A quickly performs bar and calls reset
- ▶ A then calls S.WaitOne, which closes the turnstile. But there might still be threads blocked!

Solution

- ▶ We can use a *two-phase turnstile*
- ▶ Visually:



- ▶ Two turnstiles: T1 and T2.
- ▶ Never both open (>0) at same time

Code

```
1  class Barrier{
2      private int max, num=0;
3      private object L = new object();
4      private Semaphore T1 = new Semaphore(0,1);
5      private Semaphore T2 = new Semaphore(0,1);
6      public Barrier(int max){
7          this.max=max;
8      }
9      public void wait(){
10         lock(L){
11             num++;
12             if( num == max ){
13                 for(int i=0;i<max;++i)
14                     T1.Release();
15             }
16         }
17         T1.WaitOne();
18     }
19     public void reset(){
20         lock(L){
21             num--;
22             if( num == 0 ){
23                 for(int i=0;i<max;++i)
24                     T2.Release();
25             }
26         }
27         T2.WaitOne();
28     }
29 }
```


Solution

- ▶ If you think it looks kind of like two of the simple barriers we had before: You're right!
- ▶ Can we prove correctness?
- ▶ Suppose we have our threads operating like so:

```
while(true){  
    region A  
    b.wait()  
    region B  
    b.reset()  
    region C  
}
```

Proof


- ▶ No thread can enter region B until all threads have called `b.wait()`
 - ▶ Because T1 is initially zero
 - ▶ T1 only released when all threads ready to go
- ▶ When all threads have exited `b.wait()`: T1 is zero again
- ▶ No thread can loop around to "capture" a T1 Release meant for another thread
 - ▶ Would have to pass through T2, but T2 is 0 until all threads done with region B
 - ▶ And no thread can be done with region B until it has completed `b.wait()` and seen its T1 Release
- ▶ No thread can leave region B until all threads ready to do so
 - ▶ T2 = 0 initially
 - ▶ T2 is not incremented until all threads at `b.reset()`
 - ▶ T2 = 0 when all threads complete `b.reset()`
- ▶ No thread can loop around to grab a T2.release meant for another thread
 - ▶ Would have to pass through T1, but T1 doesn't get incremented until all threads have reached `b.wait()`

Fair Semaphores

- ▶ Recall: Semaphores give no guarantees of wake-up order
 - ▶ Implementation might use a *stack* internally
 - ▶ Last asleep is first awakened
- ▶ If enough contention, a thread might *never* wake up!

Fair Semaphore

- ▶ We can create a fair semaphore that prevents starvation
- ▶ Again, we use the two-phase turnstile approach
- ▶ Visually:



- ▶ A = Not in critical section and not contending for CS
- ▶ B = Contending for CS
- ▶ C = In CS

Fair Lock

- ▶ Initially: T1 open, T2 closed
- ▶ One or more threads arrive at T1
 - ▶ Maybe simultaneously
- ▶ T1 closes
 - ▶ New threads will stack up at T1, waiting to enter
- ▶ T2 opens for one thread to enter CS
- ▶ When thread done with CS: It opens T2 for another thread
- ▶ When all threads drained from B: Close T2, open T1
- ▶ Guarantee: When thread *t* does `acquire()`, only finitely many threads may successfully enter CS before *t* enters CS

Fair Lock

```
1 class FairLock{
2     private Semaphore T1 = new Semaphore(0,1), T2 = new Semaphore(0,1);
3     private object L = new object();
4     private int numA=0, numB=0;
5     public void acquire(){
6         lock(L){
7             numA++;
8         }
9         T1.WaitOne();
10        lock(L){
11            numB++;
12            numA--;
13            if( numA == 0 )
14                T2.Release();
15            else
16                T1.Release();
17        }
18        T2.WaitOne();
19        lock(L){
20            numB--;
21        }
22    }
23    public void release(){
24        lock(L){
25            if( numB == 0 )
26                T1.Release();
27            else
28                T2.Release();
29        }
30    }
```

Explanation

- ▶ L guards numA and numB
 - ▶ We always lock so visibility of updates is guaranteed
- ▶ numA counts how many threads are in region A
- ▶ If several arrive at once: They will all increment numA
- ▶ Once any thread manages to acquire T1: Will move to region B
 - ▶ Then each one to enter B allows one more waiting thread to enter B
 - ▶ Last one in opens T2 while leaving T1 shut
 - ▶ Threads file through CS one at a time
 - ▶ As threads leave CS: Each one opens T2 for next thread
 - ▶ Last thread to leave CS leaves T2 shut and re-opens T1

Sources

- ▶ B. Goetz et al. Java Concurrency in Practice. Addison Wesley.
- ▶ A. Downey. The Little Book of Semaphores.
<http://greenteapress.com/>
- ▶ M. Herlihy and N. Shvavit. The Art of Multiprocessor Programming

Created using L^AT_EX.

Main font: Gentium Book Basic, by Victor Gaultney. See <http://software.sil.org/gentium/>

Monospace font: Source Code Pro, by Paul D. Hunt. See <https://fonts.google.com/specimen/Source+Code+Pro> and <http://sourceforge.net/adobe>

Icons by Ulisse Perusin, Steven Garrity, Lapo Calamandrei, Ryan Collier, Rodney Dawes, Andreas Nilsson, Tuomas Kuosmanen, Garrett LeSage, and Jakub Steiner. See <http://tango-project.org>